

AWS Global Infrastructure :-

globally distributed hardware and datacenters

physically networked together to act as one

large resource for end customer.

Availability Zone :- Physical location made up of one or more datacenters

Each Region generally has 3 Availability Zones.


High Availability → Common Practice to run workloads in atleast 3 AZs to ensure services remain available in case one or two datacenters fail.

Subnet is associated with an availability zone

Never choose AZ,

always choose subnet associated with AZ.

AZs are within 100kms of each other)



AZs are interconnected with high bandwidth, low latency network

(dedicated metro fibre)

Fault Domain:-

Section of a network vulnerable to damage if

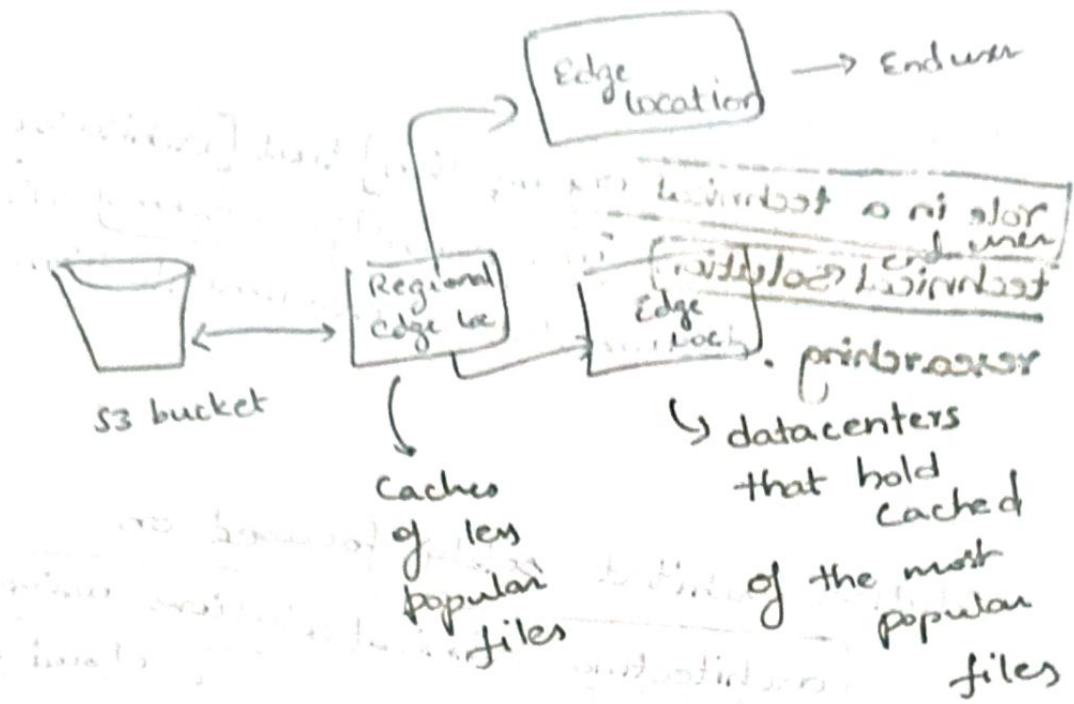
a particular device/system fails.

• Can have multiple fault domains inside fault domains

Fault level:- Collection of fault domains.

Each Amazon Region is designed to be completely isolated from other regions.
→ Achieves greatest possible fault tolerance & stability

* Each AZ is isolated but the AZs in the region are connected through low-latency links.
* Each AZ is designed as an independent failure.



Combinedly called

POP (point of presence)

AWS direct Connect:

private/ dedicated connection b/w your datacenter/office and

AWS

Direct Connect locations:-

trusted partnered data centers that

you can establish a dedicated high-speed connection to AWS.

Local Zones:-

data centers located very close to a densely populated area to provide single-digit millisecond low latency performance.